OfficeGPT RAG Chunking Test Document

This is a large paragraph designed to test how chunking behaves when the text size becomes significantly larger than the expected embedding window. The purpose of this content is to simulate real company policies or long-form documentation where multiple sentences are combined together without many line breaks. Chunking systems should ideally split this paragraph intelligently without breaking semantic meaning, while still respecting token limits required by embedding models. This is a large paragraph designed to test how chunking behaves when the text size becomes significantly larger than the expected embedding window. The purpose of this content is to simulate real company policies or long-form documentation where multiple sentences are combined together without many line breaks. Chunking systems should ideally split this paragraph intelligently without breaking semantic meaning, while still respecting token limits required by embedding models. This is a large paragraph designed to test how chunking behaves when the text size becomes significantly larger than the expected embedding window. The purpose of this content is to simulate real company policies or long-form documentation where multiple sentences are combined together without many line breaks. Chunking systems should ideally split this paragraph intelligently without breaking semantic meaning, while still respecting token limits required by embedding models.

Small paragraph one. Short and concise.
Another short paragraph for testing paragraph-based chunking.

Employee Name	Role	Department	Experience (Years)
John Doe	Developer	IT	3
Jane Smith	Designer	UI/UX	5
Arjun Patel	Manager	Operations	8

2023 2024 2025 Revenue 12000 15000 18000 Q1 3000 Q2 3500 Q3 4000 Q4 4500

This is a medium sized paragraph. It should ideally remain as a single chunk if the hybrid chunking strategy is working correctly. The system should detect that this paragraph is neither too small nor excessively large.